IZEL MEJIA

Los Angeles, CA | (661) 405-7447 | izelmejia5@gmail.com | linkedin.com/in/izel-mejia-567a67378 | github.com/izel5 | izel5.github.io

EDUCATION

University of California, Los Angeles — Los Angeles, CA

B.S. Statistics and Data Science — Expected Jun 2028

- Prior degree: A.S. Mathematics, Summa Cum Laude (Jun 2026)
- Relevant coursework: Probability Theory, Statistical Programming in R, Intro to Data Science, Data Structures, Discrete Structures, Linear Algebra, Differential Equations, Multivariable Calculus

EXPERIENCE

AGBO — Los Angeles, CA

Innovation Intern — Sep 2025 – Jun 2026

- Built the strategy team's first audience sentiment source, previously limited to sales and ratings: Playwright crawler over Reddit and web articles, comment-level chunking, content-hash deduplication, local embeddings, and LanceDB vector search behind a FastAPI + Streamlit interface. Lead engineers adopted and extended it.
- Automated a producer's multi-hour transcript reporting process into a one-click pipeline that parsed meeting transcripts, updated the project database, and generated change reports for writers and executives.
- Developed a retrieval-based review that checked new story and design ideas against confirmed project data and reported contradictions to leadership, replacing manual review that broke down on large projects.
- Designed a reference database linking design specifications and story context to concept art, used by artists to onboard to projects without digging through boards or asking colleagues.
- Cleaned and structured a ~9,000-pair bilingual lexicon from an academic source into a training-ready dataset.

PROJECTS

Consumer Review Sentiment Dashboard | *Python*, *scikit-learn*, *Streamlit*

- Cleaned and deduplicated the IMDB 50k review benchmark, then trained a TF-IDF + logistic regression classifier reaching 90.7% accuracy on a stratified 20% held-out set; dashboard shows the confusion matrix, model coefficients, and top terms.

Sleep Quality Prediction (NSDC at UCLA) | *Project Lead* — In progress

- Leading a four-person team on the NSDC Open Project "Can Data Predict How You'll Sleep Tonight?", owning scope, task split, and delivery from data to model to findings.

Forum Data Pipeline | *Python*, *LanceDB* — In progress

- Building an independent crawl-to-dataset pipeline that collects forum data, cleans and deduplicates it, and embeds each chunk locally for similarity search; next adding a SQL layer and automated analysis.

AI Research Digest | *Python*, *arXiv API*, *Claude API*

- Built a command-line tool that pulls recent papers from the arXiv API by topic and uses the Claude API to generate structured Markdown summaries for quick literature review.

Fitness Tracker Database | *Python*, *SQLite*

- Designed a three-table relational schema (users, workouts, exercises) in SQLite, loaded and queried with Python.

Causal Fairness Audit | *R* — Early stage

- Planned audit of a public dataset using the Causal Fairness Analysis framework (Plečko & Bareinboim), decomposing an outcome gap into direct, indirect, and spurious effects.

SKILLS

Languages: Python, SQL, R, Java, TypeScript

Statistics & ML: linear and logistic regression, classification, feature engineering, model evaluation (precision, recall)

Data: pandas, NumPy, scikit-learn, Excel, Power BI, PostgreSQL, SQLite, LanceDB, ETL, web scraping, vector search

AI & Tools: retrieval-augmented generation, Claude and Gemini APIs, FastAPI, Streamlit, Matplotlib, Git