IZEL MEJIA

Los Angeles, CA | (661) 405-7447 | izelmejia5@gmail.com | linkedin.com/in/izel-mejia-567a67378 | github.com/izel5 | izel5.github.io

EDUCATION

University of California, Los Angeles — Los Angeles, CA

B.S. Statistics and Data Science — Expected Jun 2028

- Prior degree: A.S. Mathematics, Summa Cum Laude (Jun 2026)
- Relevant coursework: Probability Theory, Statistical Programming in R, Intro to Data Science, Data Structures, Discrete Structures, Linear Algebra, Differential Equations, Multivariable Calculus

EXPERIENCE

AGBO — Los Angeles, CA

Innovation Intern — Sep 2025 – Jun 2026

- Built the strategy team's first audience sentiment source, previously limited to sales and ratings: Playwright crawler over Reddit and web articles, comment-level chunking, content-hash deduplication, local embeddings, and LanceDB vector search behind a FastAPI + Streamlit interface. Lead engineers adopted and extended it.
- Automated a producer's multi-hour transcript reporting process into a one-click pipeline that parsed meeting transcripts, updated the project database, and generated change reports for writers and executives.
- Developed a retrieval-based review that checked new story and design ideas against confirmed project data and reported contradictions to leadership, replacing manual review that broke down on large projects.
- Designed a reference database linking design specifications and story context to concept art, used by artists to onboard to projects without digging through boards or asking colleagues.
- Cleaned and structured a ~9,000-pair bilingual lexicon from an academic source into a training-ready dataset.

PROJECTS

Consumer Review Sentiment Dashboard | *Python*, *scikit-learn*, *Streamlit*

- Cleaned and deduplicated the IMDB 50k review benchmark, then trained a TF-IDF + logistic regression classifier reaching 90.7% accuracy on a stratified 20% held-out set; dashboard shows the confusion matrix, model coefficients, and top terms.

Sleep Quality Prediction (NSDC at UCLA) | *Project Lead* — In progress

- Leading a four-person team on the NSDC Open Project "Can Data Predict How You'll Sleep Tonight?", owning scope, task split, and delivery from data to model to findings.

College Value Analysis (STATS 20 DataFest) | *R*, *tidycensus*, *dplyr*, *ggplot2* — In progress

- Group project on whether college pays off: pulled California ACS 2022 PUMS microdata through the Census API, with 15 variables including earnings, educational attainment, field of degree, hours worked, and poverty ratio, to compare earnings across education levels and majors.

S&P 500 Market Trends (NSDC at UCLA) | *Python*, *pandas* — In progress

- Team analysis of S&P 500 and individual stock daily movements across 2,128 trading days (2012–2021), contributing an individual analysis module to the shared GitHub repo.

Forum Data Pipeline | *Python*, *LanceDB* — In progress

- Building an independent crawl-to-dataset pipeline that collects forum data, cleans and deduplicates it, and embeds each chunk locally for similarity search; next adding a SQL layer and automated analysis.

Causal Fairness Audit | *R* — Early stage

- Planned audit of a public dataset using the Causal Fairness Analysis framework (Plečko & Bareinboim), decomposing an outcome gap into direct, indirect, and spurious effects.

SKILLS

Languages: Python, SQL, R, Java, TypeScript

Statistics & ML: linear and logistic regression, classification, feature engineering, model evaluation (precision, recall)

Data: pandas, NumPy, scikit-learn, Excel, Power BI, PostgreSQL, SQLite, LanceDB, ETL, web scraping, vector search

AI & Tools: retrieval-augmented generation, Claude and Gemini APIs, FastAPI, Streamlit, Matplotlib, Git